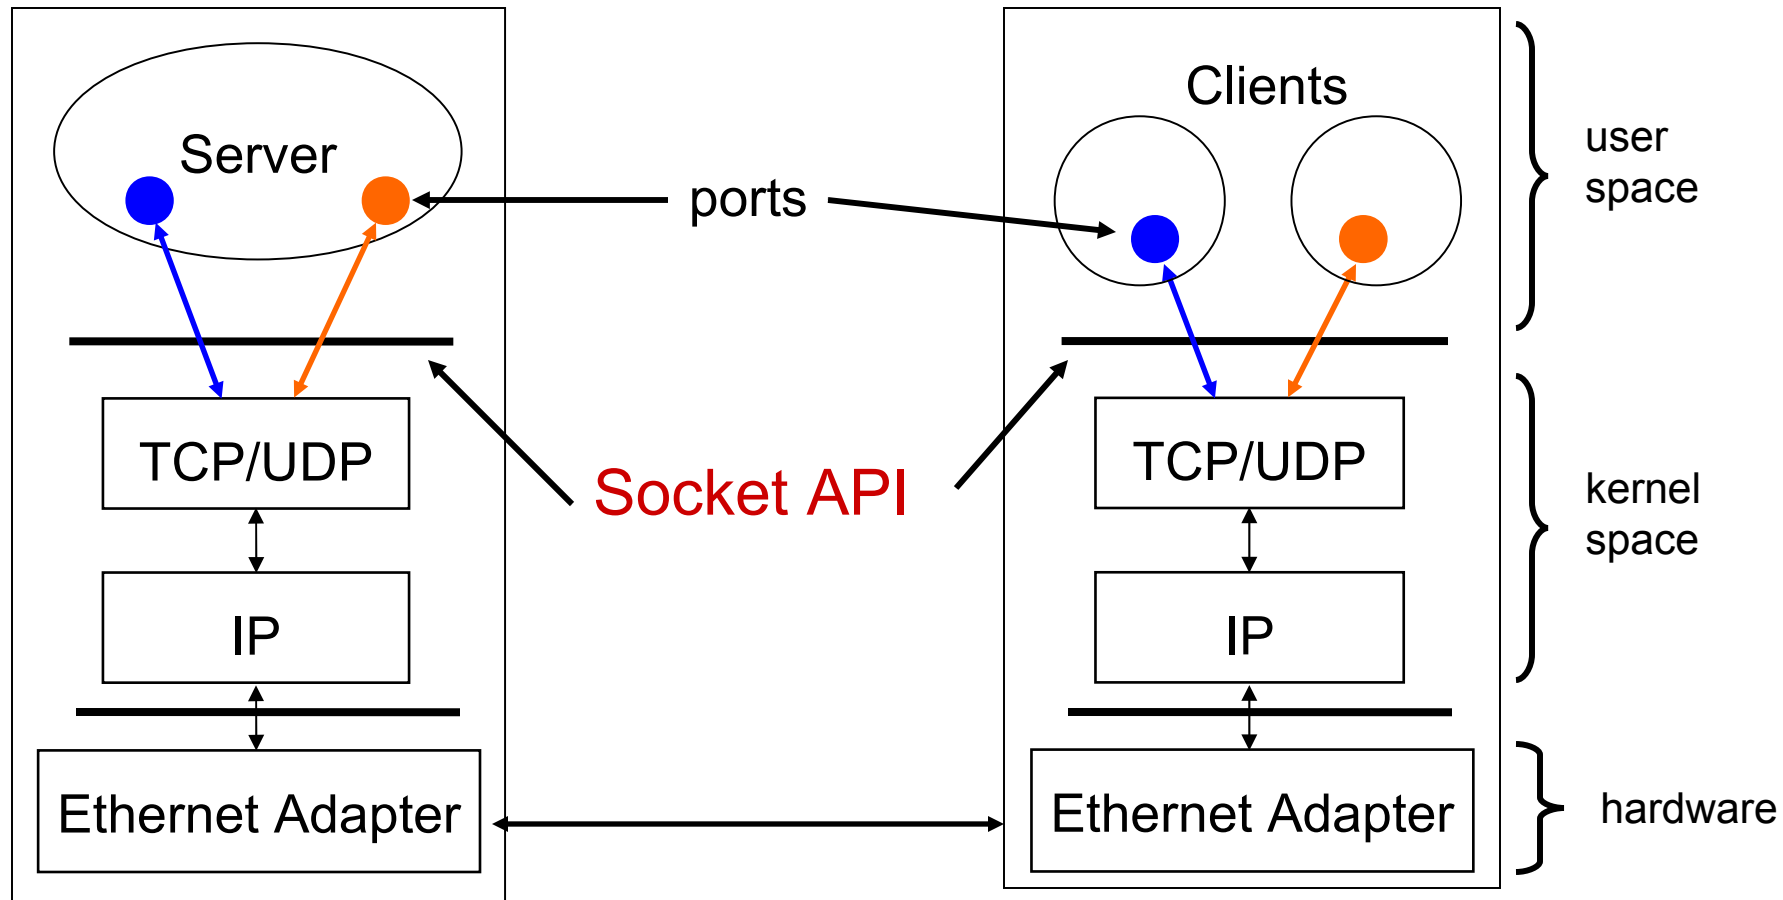


Computer Networks Lab

Socket programming in C and Java

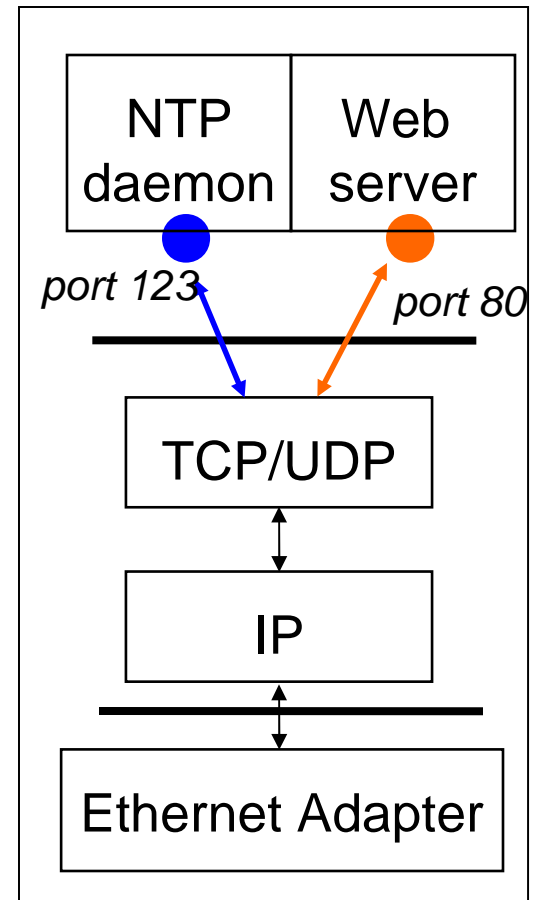
Server and Client

Server and Client exchange messages over the network through a common **Socket API**



Concept of Port Numbers

- Port numbers are used to identify “entities” on a host (16 bit)
- Port numbers can be
 - Well-known (port 0-1023)
 - Registered (port 1024 – 49151)
 - Dynamic or private (port 49152-65535)
- Servers/daemons usually use well-known ports
 - Any client can identify the server/service
 - HTTP = 80, FTP = 21, Telnet = 23, ...
 - **/etc/service** defines well-known ports
- Clients usually use dynamic ports
 - Assigned by the kernel at run time



Names and Addresses

- Each attachment point on Internet is given unique address
 - Based on location within network – like phone numbers
- Humans prefer to deal with names not addresses
 - DNS provides mapping of name to address
 - Name based on administrative ownership of host

Internet Addressing Data Structure

```
#include <netinet/in.h>

/* Internet address structure */
struct in_addr {
    u_long s_addr;          /* 32-bit IPv4 address */
};                          /* network byte ordered */

/* Socket address, Internet style. */
struct sockaddr_in {
    u_char sin_family;      /* Address Family */
    u_short sin_port;       /* UDP or TCP Port# */
                                /* network byte ordered */
    struct in_addr sin_addr; /* Internet Address */
    char sin_zero[8];       /* unused */
};
```

- `sin_family = AF_INET` selects Internet address family

Byte Ordering

```
union {  
    u_int32_t addr; /* unsigned long int, 4  
                    bytes address */  
    unsigned char c[4];  
} un;  
/* 128.2.194.95 */  
un.addr = 0x8002c25f;  
/* un.c[0] = ? */
```

c[0] c[1] c[2] c[3]

- **Big Endian**



128	2	194	95
-----	---	-----	----

– Sun Solaris, PowerPC, ...

- **Little Endian**




95	194	2	128
----	-----	---	-----

– i386, alpha, ...


- **Network byte order = Big Endian**

- Byte Order

- Different computers may have different internal representation of 16 / 32-bit integer (called **host byte order**).
- Examples
 - Big-Endian byte order (e.g., used by Motorola 68000):



- Little-Endian byte order (e.g., used by Intel 80x86):



Byte Ordering Functions

- Converts between **host byte order** and **network byte order**
 - ‘h’ = host byte order
 - ‘n’ = network byte order
 - ‘l’ = long (4 bytes), converts IP addresses
 - ‘s’ = short (2 bytes), converts port numbers

```
#include <netinet/in.h>

unsigned long int htonl(unsigned long int hostlong);
unsigned short int htons(unsigned short int hostshort);
unsigned long int ntohl(unsigned long int netlong);
unsigned short int ntohs(unsigned short int netshort);
```



What is a Socket?

- A socket is a file descriptor that lets an application read/write data from/to the network

```
int fd;          /* socket descriptor */
if ((fd = socket(AF_INET, SOCK_STREAM, 0)) < 0) {
    perror("socket");
    exit(1);
}
```

- **socket** returns an integer (socket descriptor)
 - $fd < 0$ indicates that an error occurred
 - socket descriptors are similar to file descriptors
- **AF_INET**: associates a socket with the Internet protocol family
- **SOCK_STREAM**: selects the TCP protocol
- **SOCK_DGRAM**: selects the UDP protocol

TCP Server



- For example: web server
- **What does a *web server* need to do so that a *web client* can connect to it?**

Socket I/O: `socket()`

- Since web traffic uses TCP, the web server must create a socket of type `SOCK_STREAM`

```
int fd;                /* socket descriptor */

if ((fd = socket(AF_INET, SOCK_STREAM, 0)) < 0) {
    perror("socket");
    exit(1);
}
```

- ***socket*** returns an integer (**socket descriptor**)
 - ***fd*** < 0 indicates that an error occurred
- **AF_INET** associates a socket with the Internet protocol family
- **SOCK_STREAM** selects the TCP protocol

Socket I/O: bind()

- A *socket* can be bound to a *port*

```
int fd;                                /* socket descriptor */
struct sockaddr_in srv;                /* used by bind() */

/* create the socket */

srv.sin_family = AF_INET; /* use the Internet addr family */

srv.sin_port = htons(80); /* bind socket 'fd' to port 80*/

/* bind: a client may connect to any of server addresses */
srv.sin_addr.s_addr = htonl(INADDR_ANY);

if(bind(fd, (struct sockaddr*) &srv, sizeof(srv)) < 0) {
    perror("bind"); exit(1);
}
```

- Still not quite ready to communicate with a client...

Socket I/O: listen()

- *listen* indicates that the server will accept a connection

```
int fd;                /* socket descriptor */
struct sockaddr_in srv; /* used by bind() */

/* 1) create the socket */
/* 2) bind the socket to a port */

if(listen(fd, 5) < 0) {
    perror("listen");
    exit(1);
}
```

- **Still not quite ready to communicate with a client...**

Socket I/O: accept()

- ***accept*** blocks waiting for a connection

```
int fd;                /* socket descriptor */
struct sockaddr_in srv; /* used by bind() */
struct sockaddr_in cli; /* client's add used by accept() */
int newfd;             /* returned by accept() */
int cli_len = sizeof(cli); /* used by accept() */

/* 1) create the socket */
/* 2) bind the socket to a port */
/* 3) listen on the socket */

newfd = accept(fd, (struct sockaddr*) &cli, &cli_len);
if(newfd < 0) {
    perror("accept");    exit(1);
}
```

- ***accept*** returns a new socket (***newfd***) with the same properties as the original socket (***fd***)
 - ***newfd*** < 0 indicates that an error occurred

Socket I/O: accept() continued...

```
struct sockaddr_in cli;          /* used by accept() */
int newfd;                      /* returned by accept() */
int cli_len = sizeof(cli);      /* used by accept() */

newfd = accept(fd, (struct sockaddr*) &cli, &cli_len);
if(newfd < 0) {
    perror("accept");
    exit(1);
}
```

- How does the server know which client it is?
 - **cli.sin_addr.s_addr** contains the client's *IP address*
 - **cli.sin_port** contains the client's *port number*
- Now the server can exchange data with the client by using *read* and *write* on the descriptor *newfd*.
- Why does *accept* need to return a new descriptor?

Socket I/O: read()

- *read* can be used with a socket
- ***read blocks*** waiting for data from the client but does not guarantee that `sizeof(buf)` is read


```
int fd;                /* socket descriptor */
char buf[512];         /* used by read() */
int nbytes;            /* used by read() */

/* 1) create the socket */
/* 2) bind the socket to a port */
/* 3) listen on the socket */
/* 4) accept the incoming connection */
/* 5) Create a child process to attend the connection */
/* Child will close fd, and continue communication with client */
/* parent will close newfd, and continue with accepting connection*/

if((nbytes = read(newfd, buf, sizeof(buf))) < 0) {
    perror("read"); exit(1);
}
```


TCP Client

- For example: web client
- **How does a *web client* connect to a *web server*?**



Dealing with IP Addresses

- IP Addresses are commonly written as strings (“128.2.35.50”), but programs deal with IP addresses as integers.

Converting strings to numerical address:

```
struct sockaddr_in srv;  
  
srv.sin_addr.s_addr = inet_addr("128.2.35.50") ;  
if(srv.sin_addr.s_addr == (in_addr_t) -1) {  
    fprintf(stderr, "inet_addr failed!\n"); exit(1);  
}
```

Converting a numerical address to a string:

```
struct sockaddr_in srv;  
char *t = inet_ntoa(srv.sin_addr) ;  
if(t == 0) {  
    fprintf(stderr, "inet_ntoa failed!\n"); exit(1);  
}
```

Translating Names to Addresses

- Gethostbyname provides interface to DNS
- Additional useful calls
 - Gethostbyaddr – returns `hostent` given `sockaddr_in`
 - Getservbyname
 - Used to get service description (typically port number)
 - Returns `servent` based on name

```
#include <netdb.h>

struct hostent *hp; /*ptr to host info for remote*/
struct sockaddr_in peeraddr;
char *name = "www.tezu.ernet.in";

peeraddr.sin_family = AF_INET;
hp = gethostbyname(name)
peeraddr.sin_addr.s_addr = ((struct in_addr*)(hp->h_addr))->s_addr;
```

Socket I/O: connect()

- ***connect*** allows a client to connect to a server...

```
int fd;                                /* socket descriptor */
struct sockaddr_in srv;                /* used by connect() */

/* create the socket */

/* connect: use the Internet address family */
srv.sin_family = AF_INET;

/* connect: socket 'fd' to port 80 */
srv.sin_port = htons(80);

/* connect: connect to IP Address "202.141.129.22" */
srv.sin_addr.s_addr = inet_addr("202.141.129.22");

if(connect(fd, (struct sockaddr*) &srv, sizeof(srv)) < 0) {
    perror("connect"); exit(1);
}
```

Socket I/O: write()

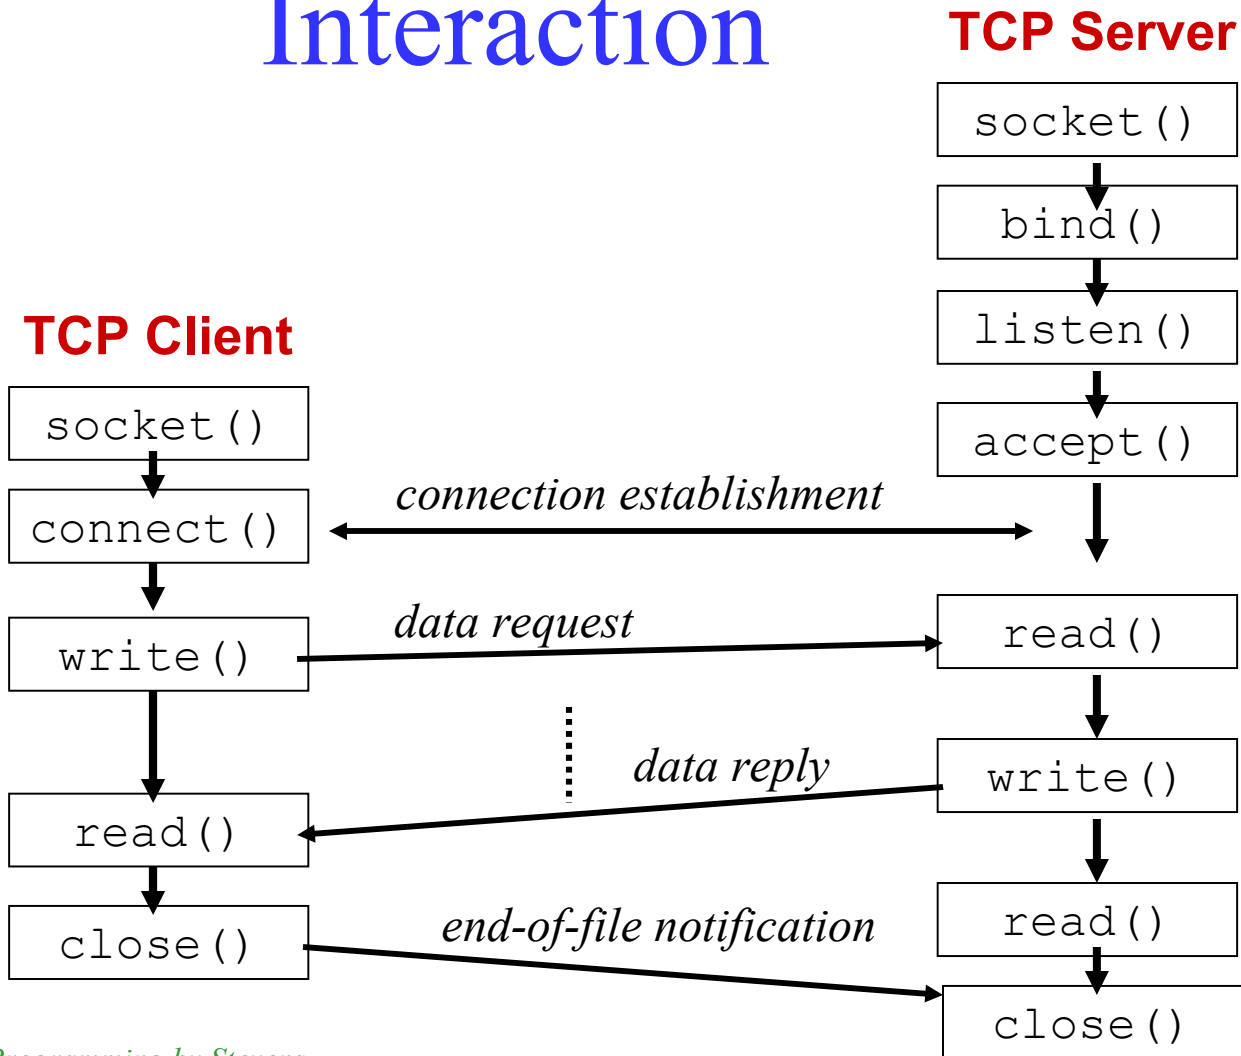
- ***write*** can be used with a socket

```
int fd;                                /* socket descriptor */
struct sockaddr_in srv;                /* used by connect() */
char buf[512];                        /* used by write() */
int nbytes;                           /* used by write() */

/* 1) create the socket */
/* 2) connect() to the server */

/* Example: A client could "write" a request to a server */
if((nbytes = write(fd, buf, sizeof(buf))) < 0) {
    perror("write");
    exit(1);
}
```

Review: TCP Client-Server Interaction



Socket programming *with TCP*

Client must contact server

- server process must first be running
- server must have created socket (door) that welcomes client's contact
- When contacted by client, **server TCP creates new socket** for server process to communicate with client
 - allows server to talk with multiple clients
 - source port numbers used to distinguish clients.

Client contacts server by:

- creating client-local TCP socket
- specifying IP address, port number of server process
- When **client creates socket**: client TCP establishes connection to server TCP

application viewpoint

TCP provides reliable, in-order transfer of bytes ("pipe") between client and server

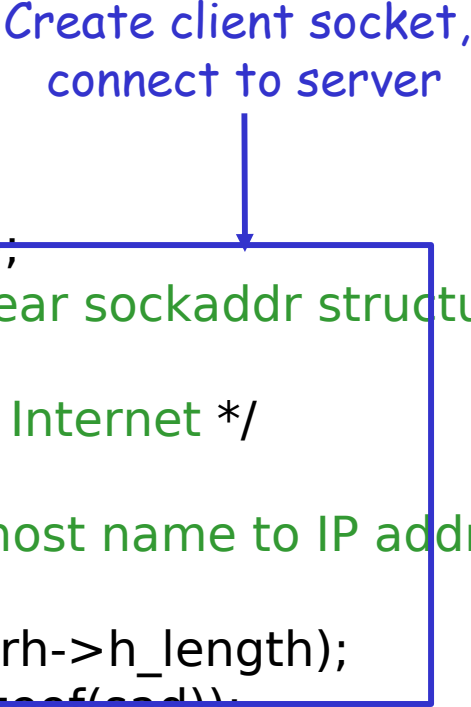
Example: C client (TCP)

```
/* client.c */
void main(int argc, char *argv[])
{
    struct sockaddr_in sad; /* structure to hold an IP address */
    int clientSocket; /* socket descriptor */
    struct hostent *ptrh; /* pointer to a host table entry */

    char Sentence[128];
    char modifiedSentence[128];

    host = argv[1]; port = atoi(argv[2]);

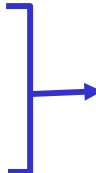
    clientSocket = socket(PF_INET, SOCK_STREAM, 0);
    memset((char *)&sad, 0, sizeof(sad)); /* clear sockaddr structure */
    /*
        sad.sin_family = AF_INET; /* set family to Internet */
        sad.sin_port = htons((u_short)port);
        ptrh = gethostbyname(host); /* Convert host name to IP address */
    */
    memcpy(&sad.sin_addr, ptrh->h_addr, ptrh->h_length);
    connect(clientSocket, (struct sockaddr *)&sad, sizeof(sad));
```



Create client socket,
connect to server

Example: C client (TCP), cont.

Get
input stream
from user



```
gets(Sentence);
```

Send line
to server



```
n=write(clientSocket, Sentence, strlen(Sentence)+1);
```

Read line
from server



```
n=read(clientSocket, modifiedSentence,  
sizeof(modifiedSentence));
```

```
printf("FROM SERVER: %s\n",modifiedSentence);
```

Close
connection



```
close(clientSocket);
```

```
}
```

Example: C server (TCP)

```
/* server.c */  
void main(int argc, char *argv[])  
{  
    struct sockaddr_in sad; /* structure to hold an IP address */  
    struct sockaddr_in cad;  
    int welcomeSocket, connectionSocket; /* socket descriptor */  
    struct hostent *ptrh; /* pointer to a host table entry */
```

```
    char clientSentence[128];  
    char capitalizedSentence[128];
```

```
    port = atoi(argv[1]);
```

Create welcoming socket at port
&
Bind a local address



```
    welcomeSocket = socket(PF_INET, SOCK_STREAM, 0);  
    memset((char *)&sad, 0, sizeof(sad)); /* clear sockaddr structure */  
    sad.sin_family = AF_INET; /* set family to Internet */  
    sad.sin_addr.s_addr = INADDR_ANY; /* set the local IP address */  
    sad.sin_port = htons((u_short)port); /* set the port number */  
    bind(welcomeSocket, (struct sockaddr *)&sad, sizeof(sad));
```

Example: C server (TCP), cont

/ Specify the maximum number of clients that can be queued */*

```
listen(welcomeSocket, 10)
```

```
while(1) {
```

Wait, on welcoming socket
for contact by a client



```
    connectionSocket=accept(welcomeSocket, (struct sockaddr *)&cad,  
    &alen);
```

```
    n=read(connectionSocket, clientSentence, sizeof(clientSentence));
```

/ capitalize Sentence and store the result in capitalizedSentence */*

```
    n=write(connectionSocket, capitalizedSentence,  
    strlen(capitalizedSentence)+1);
```

Write out the result to socket

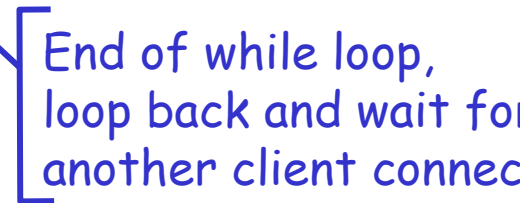


```
    close(connectionSocket);
```

```
}
```

```
}
```

End of while loop,
loop back and wait for
another client connection



Socket programming *with UDP*

UDP: no “connection”
between client and server

- no handshaking
- sender explicitly attaches IP address and port of destination to each packet
- server must extract IP address, port of sender from received packet

UDP: transmitted data may
be received out of order,
or lost

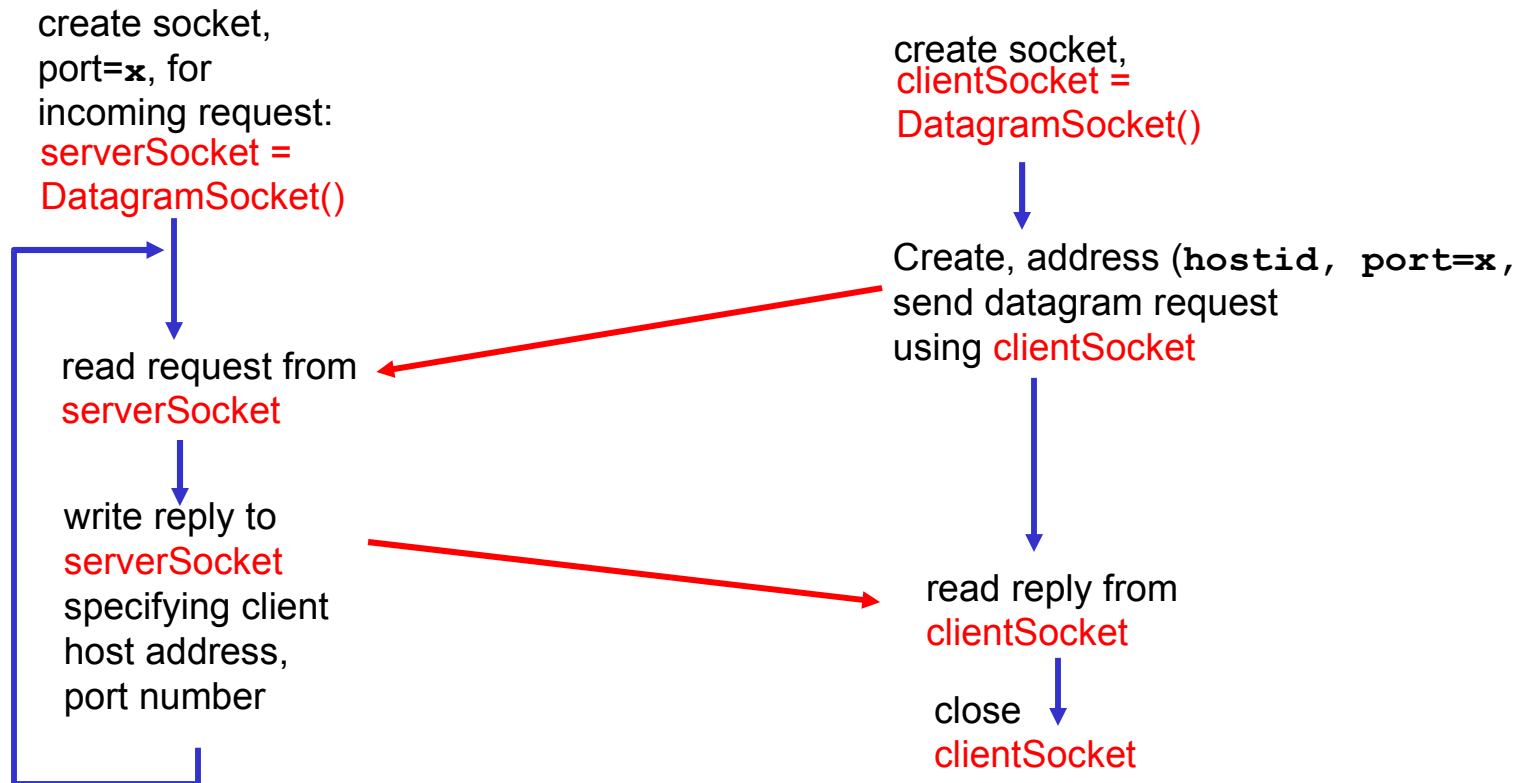
application viewpoint

*UDP provides unreliable transfer
of groups of bytes (“datagrams”)
between client and server*

Client/server socket interaction: UDP

Server (running on `hostid`)

Client



Example: C client (UDP)

```
/* client.c */
```

```
void main(int argc, char *argv[])
```

```
{
```

```
struct sockaddr_in sad; /* structure to hold an IP address */
```

```
int clientSocket; /* socket descriptor */
```

```
struct hostent *ptrh; /* pointer to a host table entry */
```

```
char Sentence[128];
```

```
char modifiedSentence[128];
```

```
host = argv[1]; port = atoi(argv[2]);
```

```
clientSocket = socket(PF_INET, SOCK_DGRAM, 0);
```

```
/* determine the server's address */
```

```
    memset((char *)&sad, 0, sizeof(sad)); /* clear sockaddr structure */
```


```
    sad.sin_family = AF_INET; /* set family to Internet */
```

```
    sad.sin_port = htons((u_short)port);
```

```
    ptrh = gethostbyname(host); /* Convert host name to IP address
```

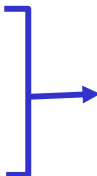
```
*/  
    memcpy(&sad.sin_addr, ptrh->h_addr, ptrh->h_length);
```

Create client socket,
NO connection to server



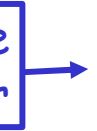
Example: C client (UDP), cont.

Get
input stream
from user




```
gets(Sentence);
```

Send line
to server



```
addr_len = sizeof(struct sockaddr);  
n=sendto(clientSocket, Sentence, strlen(Sentence)+1,  
         (struct sockaddr *) &sad, addr_len);
```


Read line
from server



```
n=recvfrom(clientSocket, modifiedSentence,  
           sizeof(modifiedSentence),  
           (struct sockaddr *) &sad, &addr_len);
```

```
printf("FROM SERVER: %s\n",modifiedSentence);
```

Close
connection



```
close(clientSocket);  
}
```

Example: C server (UDP)

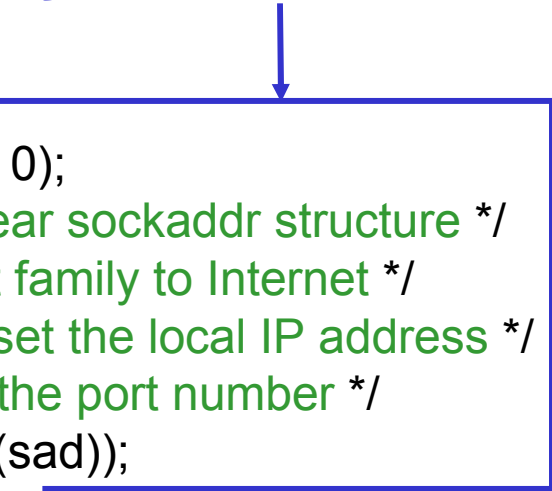
```
/* server.c */
void main(int argc, char *argv[])
{
    struct sockaddr_in sad; /* structure to hold an IP address */
    struct sockaddr_in cad;
    int serverSocket; /* socket descriptor */
    struct hostent *ptrh; /* pointer to a host table entry */

    char clientSentence[128];
    char capitalizedSentence[128];

    port = atoi(argv[1]);

    serverSocket = socket(PF_INET, SOCK_DGRAM, 0);
        memset((char *)&sad, 0, sizeof(sad)); /* clear sockaddr structure */
        sad.sin_family = AF_INET; /* set family to Internet */
        sad.sin_addr.s_addr = INADDR_ANY; /* set the local IP address */
        sad.sin_port = htons((u_short)port); /* set the port number */
    bind(serverSocket, (struct sockaddr *)&sad, sizeof(sad));
```

Create welcoming socket at port
&
Bind a local address



Example: C server (UDP), cont

```
while(1) {
```

Receive messages from clients

```
n=recvfrom(serverSocket, clientSentence, sizeof(clientSentence), 0  
            (struct sockaddr *) &cad, &addr_len );
```

```
/* capitalize Sentence and store the result in capitalizedSentence*/
```

```
n=sendto(connectionSocket, capitalizedSentence,  
strlen(capitalizedSentence)+1,0  
         (struct sockaddr *) &cad, &addr_len);
```

Write out the result to socket

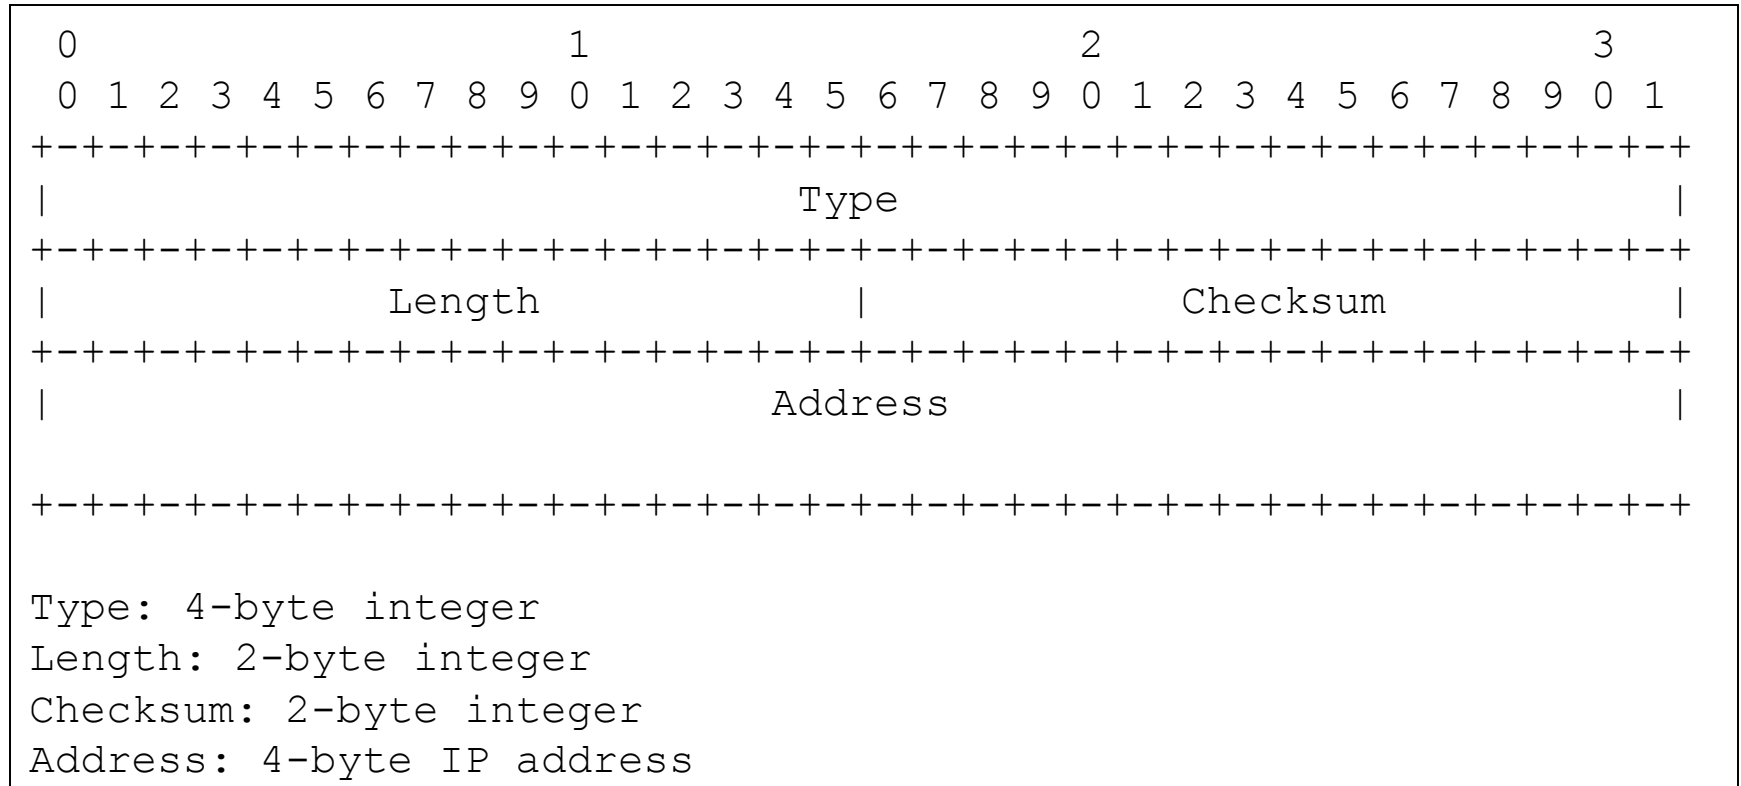
```
close(connectionSocket);
```

```
}
```

```
}
```

End of while loop,
loop back and wait for
another client connection

A Few Programming Notes: Representing Packets



A Few Programming Notes: Building a Packet in a Buffer

```
struct packet {
    u_int32_t type;
    u_int16_t length;
    u_int16_t checksum;
    u_int32_t address;
};

/* ===== */
char buf[1024];
struct packet *pkt;

pkt = (struct packet*) buf;
pkt->type = htonl(1);
pkt->length = htons(2);
pkt->checksum = htons(3);
pkt->address = htonl(4);
```

Socket Programming References

- Man page
 - usage: `man <function name>`
- Textbook
 - Stevens, Unix Network Programming, PHI


Socket Programming in Java

Sockets

- Sockets provide an interface for programming networks at the transport layer.
- Network communication using Sockets is very much similar to performing file I/O
 - In fact, socket handle is treated like file handle.
 - The streams used in file I/O operation are also applicable to socket-based I/O
- Socket-based communication is programming language independent.
 - That means, a socket program written in Java language can also communicate to a program written in Java or non-Java socket program.


Socket Communication

- A server (program) runs on a specific computer and has a socket that is bound to a specific port. The server waits and listens to the socket for a client to make a connection request.



Socket Communication


- If everything goes well, the server accepts the connection. Upon acceptance, the server gets a new socket bounds to a different port. It needs a new socket (consequently a different port number) so that it can continue to listen to the original socket for connection requests while serving the connected client.



Sockets and Java Socket Classes

- A socket is an endpoint of a two-way communication link between two programs running on the network.
- A socket is bound to a port number so that the TCP layer can identify the application that data destined to be sent.
- Java's .net package provides two classes:
 - **Socket** – for implementing a client
 - **ServerSocket** – for implementing a server

Java Sockets



It can be host_name like "agnee.tezu.ernet.in"

Implementing a Server

1. Open the Server Socket:

```
ServerSocket server;  
DataOutputStream os;  
DataInputStream is;  
server = new ServerSocket( PORT );
```

2. Wait for the Client Request:

```
Socket client = server.accept();
```

3. Create I/O streams for communicating to the client

```
is = new DataInputStream( client.getInputStream() );  
os = new DataOutputStream( client.getOutputStream() );
```

4. Perform communication with client

```
Receive from client: String line = is.readLine();  
Send to client: os.writeBytes("Hello\n");
```

5. Close sockets: client.close();

For multithreaded server:

```
while(true) {  
    i. wait for client requests (step 2 above)  
    ii. create a thread with “client” socket as parameter (the thread creates streams (as in step (3) and does  
        communication as stated in (4). Remove thread once service is provided.  
}
```

Implementing a Client

1. Create a Socket Object:

```
client = new Socket(server, port_id );
```

2. Create I/O streams for communicating with the server.

```
is = new DataInputStream(client.getInputStream() );  
os = new DataOutputStream( client.getOutputStream() );
```

3. Perform I/O or communication with the server:

- Receive data from the server:

```
String line = is.readLine();
```

- Send data to the server:

```
os.writeBytes("Hello\n");
```

4. Close the socket when done:

```
client.close();
```

A simple server (simplified code)

```
// SimpleServer.java: a simple server program
import java.net.*;
import java.io.*;
public class SimpleServer {
    public static void main(String args[]) throws IOException {
        // Register service on port 1234
        ServerSocket s = new ServerSocket(1234);
        Socket s1=s.accept(); // Wait and accept a connection
        // Get a communication stream associated with the socket
        OutputStream slout = s1.getOutputStream();
        DataOutputStream dos = new DataOutputStream (slout);
        // Send a string!
        dos.writeUTF("Hi there");
        // Close the connection, but not the server socket
        dos.close();
        slout.close();
        s1.close();
    }
}
```

A simple client (simplified code)

```
// SimpleClient.java: a simple client program
import java.net.*;
import java.io.*;
public class SimpleClient {
    public static void main(String args[]) throws IOException {
        // Open your connection to a server, at port 1234
        Socket s1 = new Socket("mundroo.cs.mu.oz.au",1234);
        // Get an input file handle from the socket and read the input
        InputStream s1In = s1.getInputStream();
        DataInputStream dis = new DataInputStream(s1In);
        String st = new String (dis.readUTF());
        System.out.println(st);
        // When done, just close the connection and exit
        dis.close();
        s1In.close();
        s1.close();
    }
}
```

Run

- Run Server on agnee.tezu.ernet.in
 - `java SimpleServer &`
- Run Client on any machine (including agnee):
 - `java SimpleClient`
Hi there
- If you run client when server is not up:
 - `sockets [1:147] java SimpleClient`
Exception in thread "main" java.net.ConnectException: Connection refused
at java.net.PlainSocketImpl.socketConnect(Native Method)
at java.net.PlainSocketImpl.doConnect(PlainSocketImpl.java:320)
at java.net.PlainSocketImpl.connectToAddress(PlainSocketImpl.java:133)
at java.net.PlainSocketImpl.connect(PlainSocketImpl.java:120)
at java.net.Socket.<init>(Socket.java:273)
at java.net.Socket.<init>(Socket.java:100)
at SimpleClient.main(SimpleClient.java:6)

Socket Exceptions

```
try {  
    Socket client = new Socket(host, port);  
    handleConnection(client);  
}  
catch(UnknownHostException uhe)  
    { System.out.println("Unknown host: " + host);  
      uhe.printStackTrace();  
}  
catch(IOException ioe) {  
    System.out.println("IOException: " + ioe); ioe.printStackTrace();  
}
```


ServerSocket & Exceptions


- public **ServerSocket**(int port) throws **IOException**
 - Creates a server socket on a specified port.
 - A port of 0 creates a socket on any free port. You can use **getLocalPort()** to identify the (assigned) port on which this socket is listening.
 - The maximum queue length for incoming connection indications (a request to connect) is set to 50. If a connection indication arrives when the queue is full, the connection is refused.
- Throws:
 - **IOException** - if an I/O error occurs when opening the socket.
 - **SecurityException** - if a security manager exists and its **checkListen** method doesn't allow the operation.

Server in Loop: Always up

// SimpleServerLoop.java: a simple server program that runs forever in a single thread

```
import java.net.*;
import java.io.*;
public class SimpleServerLoop {
    public static void main(String args[]) throws IOException {
        // Register service on port 1234
        ServerSocket s = new ServerSocket(1234);
        while(true)
        {
            Socket s1=s.accept(); // Wait and accept a connection
            // Get a communication stream associated with the socket
            OutputStream s1out = s1.getOutputStream();
            DataOutputStream dos = new DataOutputStream (s1out);
            // Send a string!
            dos.writeUTF("Hi there");
            // Close the connection, but not the server socket
            dos.close();
            s1out.close();
            s1.close();
        }
    }
}
```

Multithreaded Server: For Serving Multiple Clients Concurrently



Conclusion

- Programming client/server applications in Java is fun and challenging.
- Programming socket programming in Java is much easier than doing it in other languages such as C.
- Keywords:
 - Clients, servers, TCP/IP, port number, sockets, Java sockets